

LPS Update Meeting

October 31, 2022



AWS Assemblage dataset

- ~500k binaries in dataset
 - More than 5M: 1,142
 - More than 1M: 6,800
 - More than 100K: 115,320
- 62GB total size (all binaries, uncompressed)
- 500G total data w/ detailed JSON (roughly 7-9x binary size)

Subsetting data

- Lots of data
- Lots of smaller (not necessarily trivial) apps
- Solution: build subsets of data? >5MB, >1MB, >100k, full
 - Working to build each of these each time we export dataset
- Perhaps only full metadata for large apps?
- Notes:
 - Store rough metadata of binaries even if not actual binary / JSON
 - Start exporting subsets of data w/o deep JSON for “very small” repos (~<.5MB)

Dataset Reproducibility

- We are 3/4 done with reproducing our AWS dataset now
- Roughly ~5-10% of builds fail upon second build
 - Should be higher—we are looking into this now
- Out of the 407k binaries from second round:
 - 39k binaries are not in the first round (difference)
 - 368k are same as first round (successfully rebuilt)

Ddisasm Tuning

- Looking to use Assemblage data for novel binary analysis directions
- Idea: ddisasm uses hand-set *weights* to tune disassembly
 - Entrypoint identification
 - Symbolization
- Both of these are binary inferences, decided using weighted aggregate queries

Vector of weights

```
577
578 block_points(Block,"data",0,3,"relative-jump-table-start"):-
579     data_block_candidate(Block,_),
580     relative_address_start(Block,_,_,_,_).
581
582 //////////////////////////////////////
583 // code block candidate points
584
585 block_is_overlapping(Block,"code"),
586 block_points(Block,"code",0,-3,"overlaps with relocation"):-
587     (
588         binary_isa("X86");
589         binary_isa("X64")
590     ),
591     code_in_block_candidate_refined(EA,Block),
592     (
593         // Block beginning intersects relocation bytes:
594         relocation_size(Type,Size),
595         relocation(Target,Type,_,_,_,_),
596         EA >= Target, EA < Target + (Size/8)
597     );
598     // Relocation target is in block but not aligned with operand offsets.
599     instruction(EA,Size,_,_,_,_,_,_),
600     relocation(Target,_,_,_,_,_),
601     Target > EA, Target < EA + (Size/8),
602     !instruction_immediate_offset(EA,_,Target-EA,_),
603     !instruction_displacement_offset(EA,_,Target-EA,_),
604     ).
605
606 block_points(Block,"code",0,0,"basic point"):-
607     block_is_overlapping(Block,"code"),
608     code_in_block_candidate_refined(_,Block).
609
610 block_points(Block,"code",0,20,"start point"):-
611     block_is_overlapping(Block,"code"),
612     entry_point(Block).
613
614 block_points(Block,"code",0,1,"code section start"):-
615     block_is_overlapping(Block,"code"),
616     code_section(Section),
```

Ddisasm Tuning (Progress)

- Key idea: can get **ground truth** from Assemblage
 - This is the our key strategic advantage in this space
- Running ddisasm to *recompile* binaries has been hard
 - Linker issues, etc...
- May focus on Linux binaries to test feasibility of this idea while we reach out to GT surrounding rebuilding Windows exes